G-EVAL Template

You will be given one evaluation instance consisting of a question, an hallucinated agent trajectory, an expected hallucination explanation, and a generated hallucination explanation. The agent trajectory is a hallucinated attempt to answer the question, where the expected explanation is a human-annotated description of the earliest decisive cause underlying the hallucination, and the generated explanation is the evaluator model's predicted attribution.

<Question>**{question}**</Question>
<Agent Trajectory>**{trajectory}**</Agent Trajectory>
<Expected Explanation>**{expected_explanation}**</Expected Explanation>
<Generated Explanation>**{generated_explanation}**</Generated Explanation>

Your task is to evaluate the accuracy of generated explanation using the expected explanation as the gold reference.

Please make sure you read and understand following instructions carefully.

Evaluation Criteria:
Explanation Accuracy (1-5): the alignment between the generated explanation and the expected explanation in terms of error relevance, localization accuracy and causal correctness.
1. Score 1 (Fabricated): The explanation is distracted by irrelevant trajectory details and thus fails to attribute the error to the true hallucination cause, while also introducing fabricated evidence.
2. Score 2 (Mislocalized): The explanation references a genuine error in the trajectory but mislocalizes the hallucination by attributing it to a later step instead of the earliest decisive cause.
3. Score 3 (Wrong Cause at Correct Step): The explanation is grounded on the correct decisive step from the expected explanation, yet it misattributes the underlying cause of the hallucination.
4. Score 4 (Mostly Correct but Incomplete): The explanation matches the expected main cause with trajectory support, but is slightly incomplete or imprecise.
5. Score 5 (Exact Grounded): The explanation exactly matches the expected cause, is explicitly grounded in the trajectory, and adds no unsupported or contradictory content.

Evaluation Steps:
1. Read the question, agent trajectory, and expected explanation to establish the task intent and the trajectory segment where the hallucination becomes outcome-determining.
2. Indentify the key attribution claims in the generated explanation including the claimed error source, the described error mechanism, and the specific trajectory step it relies on.
3. Compare the generated explanation to the expected explanation for the specified evaluation criterion in terms of error relevance, localization accuracy and causal correctness.
4. Assign a single explanation accuracy score in <1,2,3,4,5> according to the criteria.
5. Ignore formatting mismatches and evaluate based on content-level correspondence only.

Answer the score from <1, 2, 3, 4, 5> and nothing else (in particular, never start your response with "I"):